

# Siege Time Series

Tre Hill

## Data Cleaning

```
library(tidyverse)
library(dplyr)
library(lubridate)
library(fpp3)
library(fable)
library(forecast)
library(tseries)
library(ggplot2)
library(scales)

#Loading and cleaning data

#Player data
data<-read.csv("player data.csv")
data <- data |> map_df(rev) |> rename("Avg.Players"="Avg..Players", "%.Gain"="X..Gain")

#New and mid season update data
up<-read.csv("patch updates.csv")
up <- up |> map_df(rev) |> mutate(new.season = ifelse(Note == "Operation Release",1,0),
  mid.season = ifelse(Note == "Mid-Season Reinforcements",1,0),
  Date = sub("^[-]*-", "", Date))

#Information on free week and weekend deals
weeks<-read.csv("free week.csv")
weeks<- weeks |> mutate(fw=if_else(Label=="Free Weekend",1,
  if_else(Label=="Free Week",2,0)),
  Date = sub("^[-]*-", "", Date)) |> select(-c("X30.Day.Peak"))

#Remove duplicate and missing from weeks
weeks<-weeks[-c(17,23),]

#Converting variables to proper type
data$Avg.Players<-as.numeric(gsub(',', '', data$Avg.Players)) #Avg.Player -> numeric
data$Gain<-as.numeric(gsub(',', '', data$Gain)) #Avg number Gained (first difference) -> numeric
data$`%.Gain`<-gsub(',', '', data$`%.Gain`)
data$`%.Gain`<-as.numeric(gsub('%', '', data$`%.Gain`)) #Percent Gained -> numeric

year_part <- as.integer(substr(data$Month, 1, 2)) + 2000 # e.g. "15" -> 2015
month_part <- substr(data$Month, 4, 6) # e.g. "Dec"
```

```

# Create a date string like "2015-Dec-01" and parse
date_str <- paste(year_part, month_part, "01", sep="-")
data$Month <- yearmonth(as.Date(date_str, format="%Y-%b-%d"))

up$Date<-as.Date(paste("01-", up$Date, sep = ""), format = "%d-%b-%y")
up$Date<-yearmonth(up$Date)

weeks$Date<-as.Date(paste("01-", weeks$Date, sep = ""), format = "%d-%b-%y")
weeks$Date<-yearmonth(weeks$Date)

# Join with corrected dates
data <- data %>%
  left_join(up %>% select(Date, new.season, mid.season), by = c("Month" = "Date")) %>%
  left_join(weeks %>% select(Date, fw), by = c("Month" = "Date")) %>%
  mutate(
    new.season = ifelse(is.na(new.season), 0, new.season),
    mid.season = ifelse(is.na(mid.season), 0, mid.season),
    fw = ifelse(is.na(fw), 0, fw)
  ) %>%
  group_by(Month) %>%
  filter(
    if(any(new.season == 1)) new.season == 1 else row_number() == 1
  ) %>%
  filter(
    if(any(mid.season == 1)) mid.season == 1 else row_number() == 1
  ) %>%
  ungroup()

```

#Creating tsibble

```

ts<-as_tsibble(data, index=Month)
ts <- ts |>
  mutate(
    covid = ifelse(
      between(Month, as.Date("2020-03-01"), as.Date("2023-05-31")),
      1,
      0
    )
  )

```

#Train/val/test split

```

train<-ts[c(1:92),]
val<- ts[c(93:104),]
test<- ts[-c(1:104),]

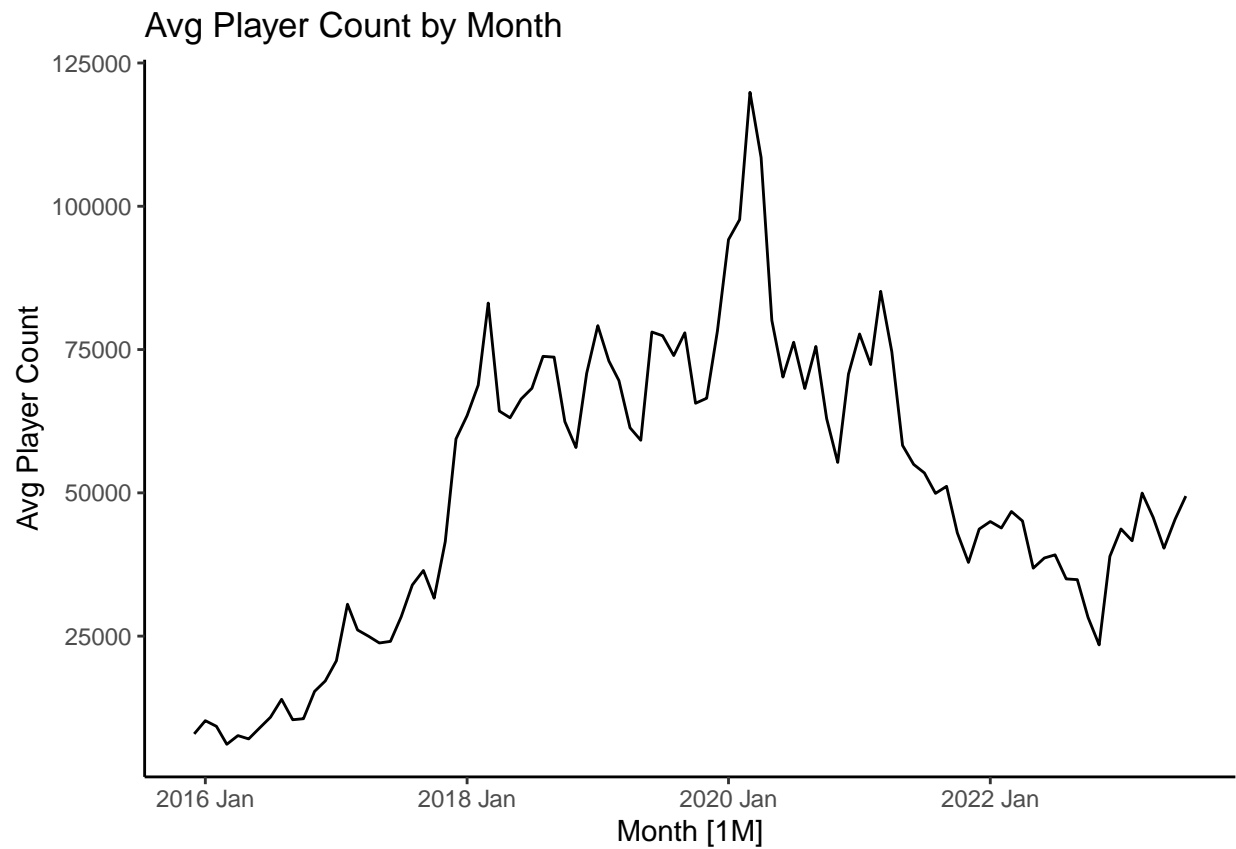
```

##EDA


```

#Time plot of training data
autoplot(train,Avg.Players) + labs(title= "Avg Player Count by Month",
  y= "Avg Player Count") + theme_classic()

```



```
#Season plot  
train |> gg_season(sqrt(Avg.Players), labels = "both") + labs(title = "Player Avg by Season",  
y = "Avg Player Count")
```

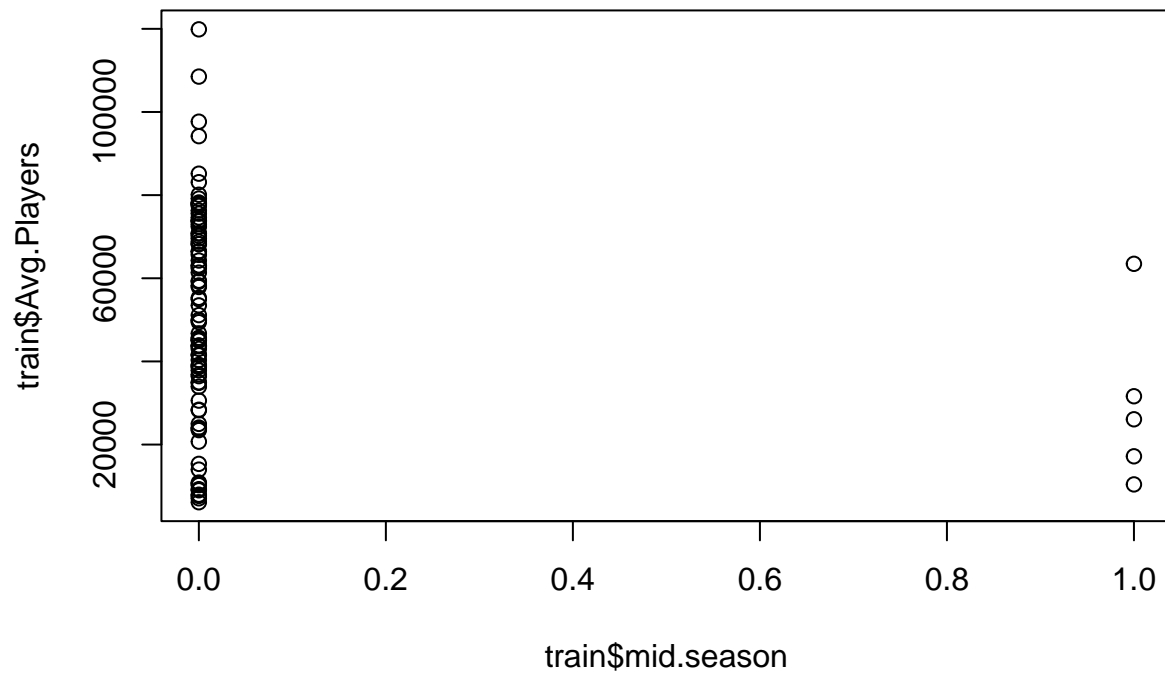


The data trends upward until COVID and then begins to drop.

Exploring possible predictors to add

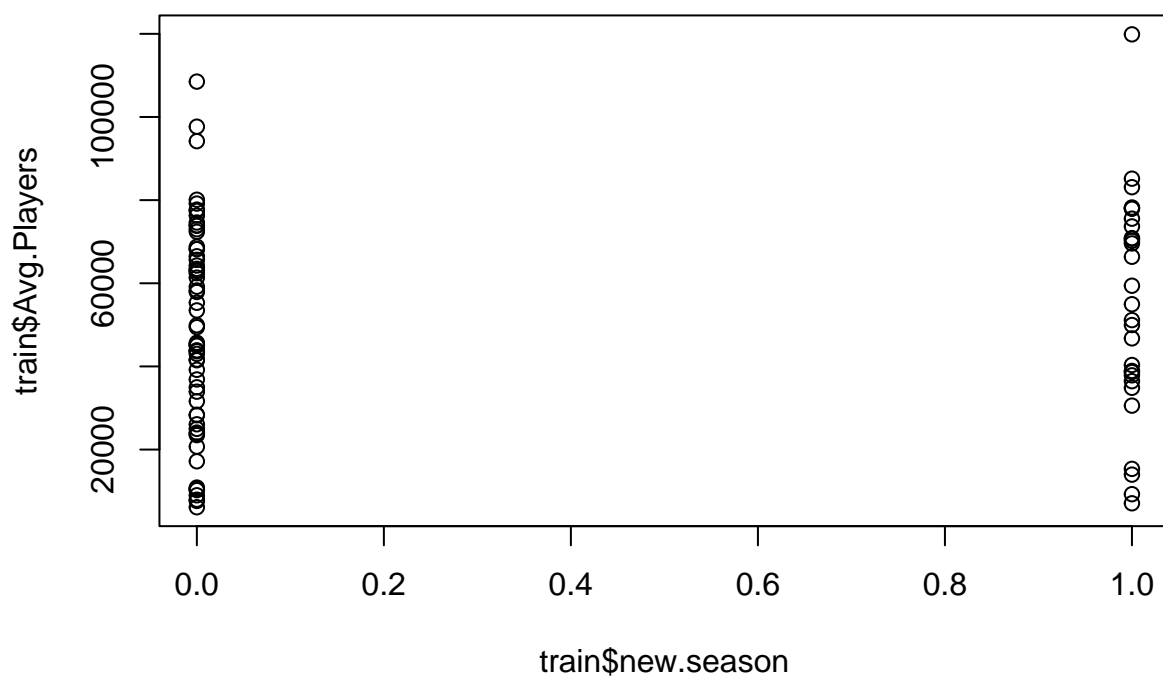
```
plot(train$mid.season, train$Avg.Players, main = "Average Players by mid season update")
```

## Average Players by mid season update




```
plot(train$new.season, train$Avg.Players, main = "Average Players by new season update")
```

## Average Players by new season update



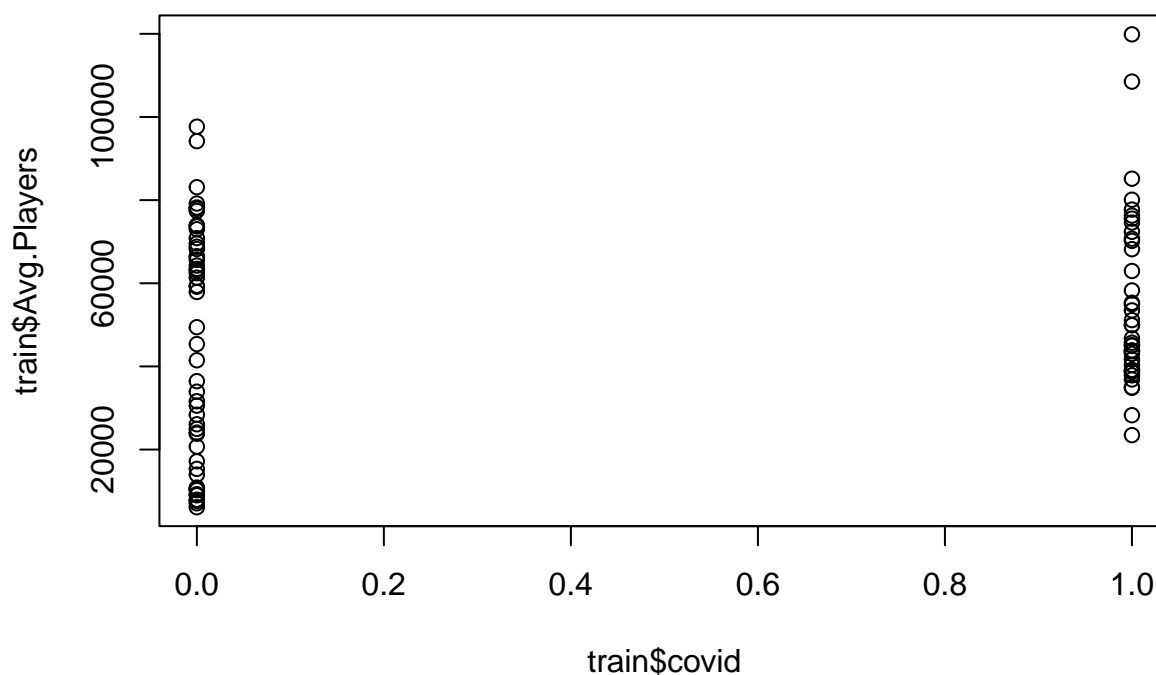
```
plot(train$fw, train$Avg.Players, main = "Average Players by free week/weekend")
```

## Average Players by free week/weekend



```
plot(train$covid, train$Avg.Players, main = "Average Players by COVID")
```

## Average Players by COVID



None of these have any separation. So, there is no point in adding these as predictors variables.


```
#Checking good box-cox
```

```
train |> features(Avg.Players, features = guerrero)
```


```
## # A tibble: 1 x 1  
##   lambda_guerrero  
##             <dbl>  
## 1             0.383
```

```
autoplot(train,sqrt(Avg.Players)) + labs(title= "Square Root of Player Count by Month",  
                                           y= "Avg Player Count") + theme_classic()
```





```
#Seasonal plot
train |> gg_season(sqrt(Avg.Players), labels = "both") + labs(title = "Player AVG by Season (sqrt)",
  y = "Avg Player Count")
```



The variances are more stable when transformed. So, modeling will focus on transformed data.

```
#STL Decomposition to estimate seasonality
dcmp<- train |> model(stl=STL(sqrt(Avg.Players)))
```


```
train |> features(sqrt(Avg.Players), feat_stl)
```

```
## # A tibble: 1 x 9
##   trend_strength seasonal_strength_year seasonal_peak_year seasonal_trough_year
##   <dbl>          <dbl>          <dbl>          <dbl>
## 1      0.972      0.645            4            0
## # i 5 more variables: spikiness <dbl>, linearity <dbl>, curvature <dbl>,
## #   stl_e_acf1 <dbl>, stl_e_acf10 <dbl>
```

```
components(dcmp) |> autoplot()
```

## STL decomposition

``sqrt(Avg.Players)` = trend + season_year + remainder`



```
dcmp[[1]] [[1]] [["fit"]] [["seasons"]] [["season_year"]] [["period"]]
```

```
## [1] 12
```

#Baseline models The models, mean naive, seasonal naive and drift will be the baseline to see if something better can be found.

```
root.baseline <- train |>
  model(
    mean.model = MEAN((Avg.Players)^(1/2)), # Forecasts mean of entire training dataset
    naive = NAIVE((Avg.Players)^(1/2)), # Forecasts last observation into future predictions
    snaive = SNAIVE((Avg.Players)^(1/2)), # Forecasts last observation of each season into future
    drift = RW((Avg.Players)^(1/2) ~ drift())
  )

#MAPE value for square root transformations
root.base.fore<- root.baseline |> fabletools::forecast(h = 12)
fabletools::accuracy(root.base.fore, val) |> arrange(MAPE)
```

```
## # A tibble: 4 x 10
##   .model   .type    ME  RMSE   MAE   MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>    <chr>  <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 drift    Test   3914. 18374. 11583. -0.434 17.4   NaN   NaN  0.457
## 2 naive    Test   8444. 20334. 12506.  7.33  17.6   NaN   NaN  0.497
```

```
## 3 mean.model Test    9184. 20937. 12957.  8.49   18.2   NaN   NaN 0.515
## 4 snaive      Test   17051. 23025. 17106. 25.1   25.2   NaN   NaN 0.447
```

Drift model performed the best. Other models will be compared to the drift model as a baseline.

##Exponential Smoothing Models

```
#Square root models
root.fit <- train |>
  model(
    SES = ETS((Avg.Players)^(1/2) ~ error("A") + trend("N") + season("N")),
    `Linear` = ETS(sqrt(Avg.Players)^(1/2) ~ error("A") + trend("A") + season("N")),
    `Damped Linear` = ETS(sqrt(Avg.Players)^(1/2) ~ error("A") + trend("Ad") +
      season("N")),
    `Damp Mult` = ETS(sqrt(Avg.Players)^(1/2) ~ error("M") + trend("Ad") +
      season("M")),
    HWAdd = ETS((Avg.Players)^(1/2) ~ error("A") + trend("A") + season("A")),
    HWMult = ETS((Avg.Players)^(1/2) ~ error("M") + trend("A") + season("M")),
    algo = ETS((Avg.Players)^(1/2)) # "algo" uses automated procedure to determine best model
  )
report(root.fit) |> arrange(AICc)
```

```
## # A tibble: 7 x 9
##   .model      sigma2 log_lik   AIC  AICc   BIC    MSE    AMSE    MAE
##   <chr>      <dbl>  <dbl> <dbl> <dbl> <dbl>  <dbl>  <dbl>  <dbl>
## 1 Damp Mult    0.00129  -138.  311.  321.  357.   0.186   0.322  0.0241
## 2 Linear       0.348   -157.  325.  325.  337.   0.332   0.665  0.469
## 3 Damped Linear 0.345   -157.  325.  326.  340.   0.327   0.635  0.465
## 4 HWAdd       194.   -442.  917.  925.  960.  161.   297.   9.80
## 5 algo        194.   -442.  917.  925.  960.  161.   297.   9.80
## 6 HWMult      0.00543  -449.  932.  940.  975.  156.   293.   0.0490
## 7 SES         296.   -469.  944.  944.  951.  290.   592.  13.6
```

Checking MAPE of the best ESM model (Damped Multiplicative)

```
root.fore <- root.fit |> select(`Damp Mult`) |> fabletools::forecast(h = 12)
fabletools::accuracy(root.fore, val) |> select(.model, MAPE)
```

```
## # A tibble: 1 x 2
##   .model    MAPE
##   <chr>    <dbl>
## 1 Damp Mult 19.3
```

##ARIMA

#ARIMA EDA


```
#Checking for ARIMA Seasonality
train |>
features(sqrt(Avg.Players), unitroot_nsdiffs) #1 seasonal difference needed
```

```
## # A tibble: 1 x 1
##   nsdiffs
##   <int>
## 1      1
```

```
train %>%
  mutate(root.seas = difference(sqrt(Avg.Players), lag = 12)) %>%
  features(root.seas, unitroot_ndiffs) #1 regular difference needed
```

```
## # A tibble: 1 x 1
##   ndiffs
##   <int>
## 1     1
```

```
train %>%
  mutate(root.seas = difference(difference(sqrt(Avg.Players), lag = 12), 1) %>% autoplot(root.seas)
```






```
#Checking if transformed data is stationary; alpha = 0.05
root_diff <- diff(diff(train$Avg.Players^(1/2), lag = 12), lag = 1)
adf.test(root_diff) #Ha: Stationary; Reject Ho-> Stationary
```

```
##
## Augmented Dickey-Fuller Test
##
## data: root_diff
## Dickey-Fuller = -3.9576, Lag order = 4, p-value = 0.01599
## alternative hypothesis: stationary
```

```
kpss.test(root_diff, null = "Level") #Ha: Non-stationary; FTR Ho-> Stationary
```

```
##
## KPSS Test for Level Stationarity
##
## data: root_diff
## KPSS Level = 0.12652, Truncation lag parameter = 3, p-value = 0.1
```

```
train %>%
  mutate(root.seas = difference(difference(sqrt(Avg.Players), lag = 12)),1) %>% gg_tsdisplay(root.seas,
```



#ARIMA modeling

Fitting 2 stochastic models: 1. with an automated procedure, and another based on the significant spikes from the ACF and pACF plots

Also exploring deterministic models with Fourier terms

```
root.arma<-train |> model( stoch=ARIMA(sqrt(Avg.Players)),
  arma11=ARIMA(sqrt(Avg.Players)~pdq(0,1,0)+PDQ(1,1,1)),
  `K = 1` = ARIMA(Avg.Players^(1/2) ~ fourier(K = 1) + PDQ(0,0,0)),
  `K = 2` = ARIMA(Avg.Players^(1/2) ~ fourier(K = 2) + PDQ(0,0,0)),
  `K = 3` = ARIMA(Avg.Players^(1/2) ~ fourier(K = 3) + PDQ(0,0,0)),
  `K = 4` = ARIMA(Avg.Players^(1/2) ~ fourier(K = 4) + PDQ(0,0,0)),
  `K = 5` = ARIMA(Avg.Players^(1/2) ~ fourier(K = 5) + PDQ(0,0,0)),
  `K = 6` = ARIMA(Avg.Players^(1/2) ~ fourier(K = 6) + PDQ(0,0,0))
root.arma |> select(stoch)
```

```
## # A mable: 1 x 1
##           stoch
##           <model>
## 1 <ARIMA(0,1,0)(0,1,1)[12]>
```

```
report(root.arima) |> arrange(AICc)
```

```
## # A tibble: 8 x 8
##   .model sigma2 log_lik   AIC   AICc   BIC ar_roots  ma_roots
##   <chr>   <dbl>   <dbl> <dbl> <dbl> <dbl> <list>   <list>
## 1 stoch    219.   -327.  657.  658.  662. <cpl [0]> <cpl [12]>
## 2 arma11   221.   -327.  659.  660.  667. <cpl [12]> <cpl [12]>
## 3 K = 2    160.   -357.  732.  734.  755. <cpl [2]> <cpl [2]>
## 4 K = 4    175.   -360.  738.  740.  760. <cpl [0]> <cpl [0]>
## 5 K = 5    179.   -360.  742.  745.  769. <cpl [0]> <cpl [0]>
## 6 K = 6    180.   -360.  743.  747.  773. <cpl [0]> <cpl [0]>
## 7 K = 3    215.   -370.  755.  756.  772. <cpl [0]> <cpl [0]>
## 8 K = 1    242.   -376.  767.  768.  784. <cpl [0]> <cpl [4]>
```

Testing best stochastic and deterministic ARIMAs against validation

```
root.arima.fore<- root.arima |> select(c(stoch,`K = 2`)) |> fabletools::forecast(h=12)
fabletools::accuracy(root.arima.fore, val) |> arrange(MAPE)
```

```
## # A tibble: 2 x 10
##   .model .type    ME  RMSE   MAE   MPE  MAPE  MASE RMSSE  ACF1
##   <chr>  <chr>   <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 stoch Test  10679. 18335. 12440.  13.4  17.3   NaN   NaN  0.475
## 2 K = 2 Test  10936. 19242. 13082.  13.4  18.2   NaN   NaN  0.495
```

##Final Test

```
tune.train<-ts[c(1:104),]

best.ar<- tune.train |> model(ARIMA(sqrt(Avg.Players)~pdq(0,1,0)+PDQ(0,1,1)))

best.fore<-best.ar |> fabletools::forecast(h = 12)
fabletools::accuracy(best.fore,test)
```

```
## # A tibble: 1 x 10
##   .model           .type    ME  RMSE   MAE   MPE  MAPE  MASE RMSSE  ACF1
##   <chr>           <chr>   <dbl> <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 ARIMA(sqrt(Avg.Players)~ Test  379. 9466. 6538. -1.25 10.3   NaN   NaN  0.382
```


Graph of the forecasted values against test data

```
graph.data<- data.frame(Month=as.Date(best.fore$Month),
                        predicted=best.fore$.mean,
                        observed=test$Avg.Players)
ggplot(data = graph.data, aes(x = Month)) +
  geom_line(aes(y = predicted, color = "Predicted"), linewidth = 1.2,alpha =.4) +
```

```

geom_point(aes(y = predicted,color = "Predicted"),size = 2.2) +
geom_line(aes(y = observed, color = "Observed"), linewidth = 1.2) +
geom_point(aes(y = observed,color = "Observed"),size = 2.2) +
scale_color_manual(values = c("Predicted" = "blue", "Observed" = "red")) +
scale_y_continuous(breaks = c(40000, 50000, 60000, 70000, 80000),
                    labels = c("40,000", "50,000", "60,000", "70,000", "80,000"))+labs(x = "Year and M
theme(axis.title.y = element_text(angle = 0,vjust = .5),
      axis.text.x = element_text(angle= 45,hjust= 1,vjust= 1,
                                  margin= margin(t= 3))) +
scale_x_date(date_breaks= "1 month", date_labels= "%b %Y")

```



#### Distributional Test Forecast

```

#Distributional forecast
index_col <- tsibble::index_var(test)
test_months <- test[[index_col]] # Extract yearmonth vector from the test set

# Plot with only test months shown on x-axis
best.ar |>
  forecast(test) |>
  autoplot(test, level = 80) +
  labs(y = "Monthly Players") +
  scale_x_yearmonth(
    breaks = test_months,
    labels = label_date("%b %Y")
  ) +


```



```

scale_y_continuous(labels = label_comma()) +
theme_classic() +
theme(
  axis.title.y = element_text(angle = 0, vjust = 0.5),
  axis.text.x = element_text(angle = 45, hjust = 1, vjust = 1)
)

```




##Future Forecast

```

fore.model<- ts |> model(ARIMA(sqrt(Avg.Players)~pdq(0,1,0)+PDQ(0,1,1)))

#Point estimate forecast
autoplot(fore.model |> forecast(h=12), level = NULL) +
  labs(y = "Monthly Players") +
  scale_x_yearmonth(
    breaks = yearmonth(c("2025 Aug", "2025 Sep", "2025 Oct", "2025 Nov",
                        "2025 Dec", "2026 Jan", "2026 Feb", "2026 Mar",
                        "2026 Apr", "2026 May", "2026 Jun", "2026 Jul")),
    labels = label_date("%b %Y")
  ) +
  scale_y_continuous(labels = label_comma()) +
  theme_classic() +
  theme(
    axis.title.y = element_text(angle = 0, vjust = 0.5),
    axis.text.x = element_text(angle = 45, hjust = 1, vjust = 1)
  )

```



```
#Distributional forecast
autoplot(fore.model |> forecast(h=12), level = 80) +
  labs(y = "Monthly Players") +
  scale_x_yearmonth(
    breaks = yearmonth(c("2025 Aug", "2025 Sep", "2025 Oct", "2025 Nov",
                        "2025 Dec", "2026 Jan", "2026 Feb", "2026 Mar",
                        "2026 Apr", "2026 May", "2026 Jun", "2026 Jul")),
    labels = label_date("%b %Y")
  ) +
  scale_y_continuous(labels = label_comma()) +
  theme_classic() +
  theme(
    axis.title.y = element_text(angle = 0, vjust = 0.5),
    axis.text.x = element_text(angle = 45, hjust = 1, vjust = 1)
  )
```

